

Got it Kowshik  — let's lay everything out **clearly** for each member: **pipeline, architecture, deliverables, and sequence of execution**. This will give you and your teammates a single source of truth.

Overall End-to-End Architecture (System View)

Flow:

DICOM ingest → Preprocessing → Detection → Segmentation → Risk classification → Explainability → Findings JSON → Reporting (Clinician + Patient PDFs) → FHIR Export → Frontend Viewer (UI).

Sequence:

1. **Teja (Data/MLOps)** prepares data + endpoints + queue.
 2. **Dinesh (ML-A)** builds detection, segmentation, risk, Grad-CAM.
 3. **Kowshik (ML-B)** builds schema, reports, FHIR, Q&A.
 4. **Rahul (Frontend)** builds offline UI to tie everything together.
-

Dinesh — ML-A (Vision Lead)

Pipeline

1. **Preprocessing** → resample (1mm), HU clip, normalize, 2.5D slice packs.
2. **Detection** → LoG blob detector (v1) / CenterNet-lite (v2).
3. **Segmentation** → Mobile-UNet / Half-UNet (lightweight).
4. **Risk Head** → deep features + radiomics → MLP classifier.
5. **Explainability** → Grad-CAM++ overlays on nodules.
6. **Integration** → run_inference(study_dir) returns findings.json, masks, saliency.

Architecture

- **Model stack:**

- Detector (classical + small CNN head)
- Segmenter (2.5D/3D UNet)
- Risk classifier (MLP with features + radiomics)
- Grad-CAM++ (XAI)

- **Frameworks:** PyTorch + MONAI, SimpleITK.

- **Inputs/Outputs:**

- Input: DICOM → preprocessed volume.
- Output: JSON + overlays.

Deliverables

- ml/preprocess.py, ml/detector.py, ml/segmenter.py, ml/risk.py, ml/xai.py, ml/infer.py.
- findings.json with nodules + risk.
- Mask PNGs + saliency PNGs.
- Metrics: Dice ≥ 0.8 , AUC ≥ 0.9 , ≤ 60 s/study latency.

Sequence

- Week 1: Preprocessing + detector v1.
 - Week 2: Train UNet seg baseline, generate masks.
 - Week 3: Risk head + Grad-CAM overlays.
 - Week 4: Integrate \rightarrow run_inference() end-to-end.
-

Kowshik — ML-B (Reporting/NLP)

Pipeline

1. **Schema & Validation** \rightarrow findings.schema.json + validator.
2. **Reporting** \rightarrow template-based PDF reports: clinician (EN) + patient (EN/HI).
3. **FHIR Export** \rightarrow DiagnosticReport + ImagingStudy JSONs.
4. **Q&A** \rightarrow rule-based engine answering from findings.json.

Architecture

- **Validator:** JSON schema + rules (size, lobe, Lung-RADS).
- **Reporter:** Markdown \rightarrow PDF (WeasyPrint/ReportLab).
- **FHIR Exporter:** Structured JSON compliant with HL7 FHIR.
- **Q&A Engine:** JSON-grounded rule-based retrieval.
- **Inputs/Outputs:**
 - Input: findings.json (from Dinesh).
 - Output: clinician.pdf, patient_en.pdf, patient_hi.pdf, diagnostic_report.json.

Deliverables

- app/validators.py, app/reporter.py, app/fhir_export.py, app/qa_engine.py.
- Templates: templates/clinician_en.md, patient_en.md, patient_hi.md.
- Sample outputs: PDFs + FHIR JSONs.
- Validator passes ≥ 10 cases, no hallucination in reports.

📌 Sequence

- Week 1: Schema draft + validator v1.
 - Week 2: Report generator v1.
 - Week 3: FHIR exporter + Q&A engine.
 - Week 4: BLEU/ROUGE eval vs reference reports + polish.
-

💻 Teja — Data & MLOps

📌 Pipeline

1. **Data Handling** → DICOM loader, save overlays, save JSON.
2. **Data Versioning** → DVC pipeline for raw → processed CTs.
3. **Queue/Orchestration** → RQ/Redis or Celery to manage pipeline jobs.
4. **Evaluation Harness** → scripts for latency, FROC, Dice.

📌 Architecture

- **Backend:** FastAPI server + RQ/Celery workers.
- **DB:** Redis/SQLite for queue states.
- **Eval Scripts:**
 - latency_bench.py
 - mini_froc.py
 - dice_smallset.py
- **Inputs/Outputs:**
 - Input: DICOM studies (10–20 proto_pack).
 - Output: metrics CSVs + stored results in /artifacts/metrics/.

📌 Deliverables

- ml/io.py, data/proto_pack/, data/dataset_card.md.
- app/queue.py (job queue).
- Metrics: latency.csv, mini_froc.csv, dice.csv.
- README with one-command run.

📌 Sequence

- Week 1: Dataset prep + proto_pack.
- Week 2: Queue skeleton + endpoints (/ingest, /analyze, /status).

- Week 3: Eval scripts (latency, Dice).
 - Week 4: Crash-safe job queue, reproducible runs.
-

Rahul — Frontend

Pipeline

1. **UI Framework** → Streamlit (offline-first).
2. **Pages:**
 - Home/Ingest → upload DICOM.
 - Analyze → POST to server, poll /status.
 - Results → axial viewer, overlay toggle, downloads.
3. **Features:** bilingual toggle, offline queue (outbox), error handling.

Architecture

- **Frontend:** Streamlit UI with pages.
- **Viewer:** Slice slider, overlay toggles (mask/saliency).
- **Downloads:** PDFs, JSONs.
- **Inputs/Outputs:**
 - Input: study_uid (after ingest).
 - Output: Visual viewer + downloads.

Deliverables

- ui/app.py with pages (Ingest, Analyze, Results).
- Responsive axial viewer (512×512).
- Toggle overlays + bilingual UI.
- Download buttons working.

Sequence

- Week 1: Skeleton UI with dummy JSON.
 - Week 2: Integrate viewer + overlay toggles.
 - Week 3: Add downloads (PDFs, JSONs).
 - Week 4: Offline queue + polish UX.
-

Sequential Execution (Team)

1. **Teja (Week 1)** → prepare proto_pack + ingest/analyze endpoints.
 2. **Dinesh (Week 1–2)** → preprocessing + baseline UNet segmentation.
 3. **Kowshik (Week 1–2)** → schema + validator + draft templates.
 4. **Rahul (Week 1)** → Streamlit skeleton with dummy JSON.
 5. **Dinesh + Teja (Week 2–3)** → integrate detector + risk head.
 6. **Kowshik (Week 3)** → FHIR export + Q&A.
 7. **Rahul (Week 3)** → integrate downloads + bilingual toggle.
 8. **Final Week** → integrate all modules → dry-run demo → freeze repo.
-

- This gives you **per-member pipeline, architecture, deliverables, and sequence**.
- 👉 Do you want me to now **generate the starter repo structure + code skeletons** (empty functions, file scaffolds) for all 4 roles so your team can directly start coding?