

# FYP - Early-1

awesome — here's a super-concrete, role-by-role prototype plan you can execute immediately. It tells each of you **what to build, how to build it, what data to use, the exact files to touch, acceptance criteria, and a day-by-day schedule** to land the end-to-end demo.

## Prototype Goal (Sep 29)

A working, offline, end-to-end **Lung-CT** demo:

- **DICOM ingest → analyze → results viewer** (masks + saliency).
- **Detector+segmenter v1, risk head v0, Grad-CAM++.**
- Outputs: **findings.json, Clinician PDF (EN), Patient PDF (EN/HI), FHIR DiagnosticReport & ImagingStudy JSON.**
- Queue-based, crash-safe,  $\leq 60\text{--}90\text{ s/study}$  on RTX 4060.
- Mini-eval: latency, 1 FROC point, a few Dice examples, screenshots.

Repo root: `lungct/` (layout below). Everyone sticks to this structure.

```
bash                                         ⌂ Copy ⌂ Edit

lungct/
app/      # APIs & services
ui/       # Streamlit (or React) UI
ml/       # models & inference
data/     # proto_pack (10-20 CTs)
eval/     # quick metrics
artifacts/ # outputs (json/pdf/fhir/overlays)
config.yaml
requirements.txt
README.md
```

## Shared setup (all)

- Environment: Python 3.10; PyTorch (CUDA), pydicom, SimpleITK or MONAI, numpy, fastapi, uvicorn, streamlit, reportlab/weasyprint, onnxruntime (CPU ok for now), rq/redis or celery, pydantic, pillow, scikit-image.
- Git: branch per role (`feat/ml-a`, `feat/ml-b`, `feat/data-mlops`, `feat/frontend`); PRs into `main`.
- Data: create a **proto\_pack** of **10–20 CT studies** (de-identified). Prefer **LIDC-IDRI** DICOMs (with XML annotations) for 10 subjects + **LNDb v4** 5 subjects (NIfTI+CSV) if you already have it. Keep total size  $\leq 3\text{ GB}$ .
- For 5–8 nodules, produce **rough masks** (slice-wise contours) so we can compute Dice in the demo. (Use 3D Slicer or draw simple masks around nodules — good enough for prototype.)

# Dinesh — ML-A (Vision Lead)

## Mission

Ship the **detector+segmenter v1**, **risk head v0**, and **Grad-CAM++ overlays**, callable via one function:

```
run_inference(study_dir) -> findings.json + overlays .
```

## What to build (files & functions)

- ml/preprocess.py
  - resample\_to\_iso(vol, spacing=1.0) → SimpleITK/MONAI.
  - hu\_clip(vol, low=-1000, high=400), normalize01(vol) .
  - make\_2p5d\_packs(vol, k=5, stride=1) → returns tiles (C=5, H, W) + z indices.
- ml/detector.py
  - **v1 candidate generator (fast):** LoG blob detector across slices or small CenterNet-lite head.
  - detect\_candidates(packs) -> [(z,y,x,score)] (keep top-K per slab; N≈50–200).
- ml/segmenter.py
  - **Mobile-UNet / Half-UNet 2.5D (single-class).**
  - segment\_cubes(cubes) -> masks (threshold→morphology→3D merge).
- ml/risk.py
  - extract\_deep\_features(cubes) (global avg pooled penultimate layer).
  - quick\_radiomics(mask, vol) (sphericity, volume, mean HU).
  - risk\_head(features+radios) -> p\_malignant (tiny MLP).
- ml/xai.py
  - gradcam\_pp(model, cube) -> heatmap (over risk head's features).
  - Compose **mask + heatmap** overlays (save as PNGs).
- ml/infer.py
  - Tie it all:

```
python
```

Copy Edit

```
def run_inference(study_dir, out_dir):
    vol = load_dicom_series(study_dir)
    vol = normalize01(hu_clip(resample_to_iso(vol)))
    packs = make_2p5d_packs(vol)
    cands = detect_candidates(packs)
    cubes = extract_cubes(vol, cands, size=64)
    masks = segment_cubes(cubes)
    nods = merge_3d_blobs(masks, cands) # (centroid, size, vol)
    feats = extract_deep_features(cubes)
    radios = [quick_radiomics(m, vol) for m in masks]
    risk = risk_head(concat(feats, radios))
    sal = [gradcam_pp(risk_head, cb) for cb in cubes]
    findings = build_findings_json(nods, risk, sal)
    save_overlays(vol, masks, sal, out_dir)
    save_json(findings, out_dir/"findings.json")
```

## Data to use

- **LIDC-IDRI** DICOMs (10 subjects). For a few nodules, read XML ROI to craft quick slice masks (or rough boxes).
- Optional: **LNDb v4** samples for extra masks.

## Minimal training plan (fast)

- Train the 2.5D **Mobile-UNet** on your rough masks (patch-based). 4–6 hours on RTX 4060; batch 16; CE/Dice loss; early stop on val Dice.
- Detector v1 can be classical (LoG blob + threshold); upgrade later.

## Acceptance criteria (for you)

- 1 study end-to-end runs in  $\leq 90$  s; produces **findings.json** + mask & saliency PNGs.
  - At least 3 cases with reasonable masks; Dice  $\geq 0.70$  on the few labeled nodules.
  - One FROC point on proto\_pack: Sens@2 FP/scan computed by Teja's scripts.
- 

# Kowshik — ML-B (Reporting/NLP)

## Mission

Owning the **findings schema**, **validators**, **RSNA clinician report**, **patient EN/HI summaries**, **PDF generator**, **grounded Q&A**, and **FHIR exporter**.

## What to build (files & functions)

- `appValidators.py`
  - `validate_findings(json)` → checks count/size/site, units (mm/mm<sup>3</sup>), Lung-RADS rules.
- `appReporter.py`
  - **Templates:**
    - `templates/clinician_en.md` (Technique, Findings, Impression, Lung-RADS, Recommendation).
    - `templates/patient_en.md`, `templates/patient_hi.md` (8th-grade, glossary).
  - `render_reports(findings.json)` -> `clinician.pdf`, `patient_en.pdf`, `patient_hi.pdf` (ReportLab/WeasyPrint).
- `appFhirExport.py`
  - `to_fhir_diagnostic_report(findings.json)` -> `diagnostic_report.json`
  - `to_fhir_imaging_study(dicom_meta)` -> `imaging_study.json`
  - Reference observations for each nodule (size, location, risk).
- `appQaEngine.py`
  - **Grounded Q&A:** given a question (e.g., "Why malignant?"), return an answer **only** from findings JSON + pointers (nodule ID, slices). (Rule-based now; LLM later.)
- `appSamples/`
  - `findings.schema.json` (final), `sample_findings.json`, `sample_fhir.json`.

## Data to use

- **LNDb v4** report tables (if available) to sanity-check phrasing.

- Otherwise: craft 10 "reference" reports for proto\_pack cases (manual) to test BLEU/ROUGE later.

## Acceptance criteria (for you)

- findings.schema.json validated with at least **10 cases**.
  - Clinician PDF (EN) and Patient PDF (EN/HI) render in <2 s each; no hallucinations (strict template).
  - FHIR DiagnosticReport & ImagingStudy JSONs produce (valid JSON; cross-referenced IDs).
  - Q&A returns **evidence-linked** answers ("Nodule N1 in RUL, 7.2 mm; features: solid, spiculation absent").
- 

# Teja — Data / MLOps

## Mission

Deliver the **data pipeline, DVC-tracked proto\_pack, job queue orchestration, and the eval/benchmark scripts**. Make the demo robust and reproducible.

## What to build (files & functions)

- ml/io.py
  - load\_dicom\_series(study\_dir) -> (vol, meta) using pydicom/SimpleITK.
  - save\_overlays(vol, masks, saliency, out\_dir) → axial PNG stacks.
  - save\_json(obj, path) .
- data/
  - proto\_pack/ (10–20 CTs).
  - dataset\_card.md (source, license, preprocessing).
  - DVC pipeline: dvc init ; track raw → processed.
- app/queue.py
  - RQ/Celery workers; job model: **Ingested → Queued → Preprocess → Detect → Segment → Risk → Explain → Report → Export → Completed**; state persisted in DB.
- eval/latency\_bench.py
  - Measure per-study time, per-stage breakdown, memory/VRAM peaks; CSV output.
- eval/mini\_froc.py
  - On proto\_pack, compute a **single FROC point**: Sens@2 FP/scan (require you to store GT centers for a handful of nodules).
- eval/dice\_smallset.py
  - Dice & 95% HD on the few hand-labeled masks.

## Data to use

- LIDC-IDRI DICOMs (10). Pull series with thin slices ( $\leq 1.5$  mm).
- Use LIDC XML to extract center coords for a few nodules (for mini-FROC ground truth).
- If LNDb v4 is available, include 3–5 NIfTI volumes as external test; otherwise skip.

## Acceptance criteria (for you)

- One-command **repo**: `python ml/infer.py --study_dir ...` runs & writes outputs.
  - Queue survives a kill/restart (job resumes).
  - Produce **latency.csv**, **mini\_froc.csv**, **dice\_smallset.csv** in `/artifacts/metrics/`.
  - `README.md` with **exact run commands** and environment pins.
- 

## Rahul — Frontend

### Mission

Deliver the **offline viewer UI**: load CT, run analysis, show overlays (mask + saliency), bilingual text, downloads (PDFs, JSON), and clear progress/error states.

### What to build (files & pages)

- Choose **Streamlit** (fastest) for the prototype.
- `ui/app.py` — pages:
  1. **Home / Ingest**: select DICOM folder/zip → POST `/ingest` → shows `study_uid`.
  2. **Analyze**: button calls POST `/analyze/{study_uid}` → poll `/status/{job_id}` → progress bar with ETA.
  3. **Results**:
    - Axial viewer (slice slider), window/level controls.
    - Toggle **Masks** and **Saliency**; click a nodule → show **size, risk, lobe**.
    - Buttons: **Download Clinician PDF**, **Patient PDF (EN/HI)**, **Findings JSON**, **FHIR JSON**.
    - **Language toggle (EN/HI)**; high-contrast theme; large buttons.
- Error handling: friendly messages, retry button; show "offline" if ABDM sync isn't reachable.
- Store-and-forward: a local folder `outbox/fhir/` with queued JSONs; show counts.

### Data to use

- Use Teja's **proto\_pack**; develop with 1–2 small studies for fast iterations.
- Use `artifacts/samples/` for UI placeholders before ML is ready.

### Acceptance criteria (for you)

- Clean flow: Ingest → Analyze → Results, **without touching the terminal**.
  - Viewer is responsive for **512–768 px** axial slices; overlays toggle without lag.
  - All 4 downloads work; language toggle flips Patient PDF labels.
- 

## Interfaces (what you all agree on)

### FastAPI endpoints (Teja implements, Rahul consumes, ML writes outputs)

- POST `/ingest` → `{study_uid}` (server writes the DICOM path mapping).
- POST `/analyze/{study_uid}` → `{job_id}`.
- GET `/status/{job_id}` → `{progress: 0–100, stage: string, eta_s: int}`.

- GET /results/{study\_uid} → {links: {findings\_json, overlays\_zip, clinician\_pdf, patient\_en\_pdf, patient\_hi\_pdf, fhir\_json}} .
- POST /export/fhir/{study\_uid} → stores JSON to outbox/fhir/ , returns path.

## Findings JSON (Kowshik owns)

```
json
Copy Edit

{
  "study_uid": "1.2.840....",
  "meta": {"series_uid": "1.2.3", "spacing_mm": [1.0,1.0,1.0]},
  "nODULES": [
    {
      "id": "N1", "lobe": "RUL", "type": "solid",
      "center_zyx": [128,192,240], "long_axis_mm": 7.2, "vol_mm3": 190,
      "risk": 0.72, "uncertain": false,
      "slices": [127,128,129], "mask_path": "masks/N1_mask.npz",
      "saliency_path": "saliency/N1_axial_128.png"
    }
  ],
  "summary": {"nodule_count": 1, "largest_mm": 7.2, "lung_rads": "3"},
  "quality": {"motion": "mild", "noise": "moderate"},
  "recommendation": "CT follow-up in 6 months"
}
```

## Day-by-day plan to hit the demo

### D-10 to D-8

- Teja: DICOM loader, resample; ingest & analyze endpoints; queue skeleton; dataset card & proto\_pack.
- Dinesh: preprocess; LoG-based detector; cube extraction; Mobile-UNet code scaffold.
- Rahul: Streamlit skeleton (pages, file pickers, progress bar), dummy result page (reads sample\_findings.json ).
- Kowshik: finalize findings.schema.json ; draft templates (clinician EN, patient EN/HI); validator v0.

### D-7 to D-6

- Dinesh: segmenter train (4–6h); merge 3D blobs; size/volume.
- Teja: wire ml/infer.py into queue; save overlays; /results returns links; start latency logger.
- Rahul: axial viewer with overlays; window/level; language toggle.
- Kowshik: PDF generation; FHIR builder v0; Q&A stub grounded in JSON.

### D-5 to D-4

- Dinesh: risk head v0; Grad-CAM++; save saliency PNGs.
- Teja: mini\_froc & dice scripts (with 3–5 GT nodules); latency CSV.
- Rahul: download buttons; error handling; "store-and-forward" outbox panel.
- Kowshik: validator v1 (count/size/site); finalize templates; sample FHIR JSON.

### D-3

- Integrate: end-to-end on 3 studies; fix timings; crash/restart resume; audit log table.
- Collect screenshots.

**D-2**

- Dry-run demo; fill metrics tables; polish PDFs (fonts for Hindi).

**D-1**

- Freeze repo; export **Prototype.pdf** (screenshots, metrics, sample JSON/PDFs).
- 

**Per-role acceptance checklist (copy into your own todo)****Dinesh (ML-A)**

- `run_inference()` outputs `findings.json`, masks, saliency for any study.
- 3 demo cases with reasonable masks; Dice  $\geq 0.70$  on labeled nodules.
- Latency end-to-end  $\leq 90$  s/study on 4060.

**Kowshik (ML-B)**

- `findings.schema.json` + validator pass on 10 cases.
- Clinician PDF (EN) & Patient PDFs (EN/HI) generated, consistent with JSON.
- FHIR DiagnosticReport & ImagingStudy JSONs saved; grounded Q&A returns evidence-linked answers.

**Teja (Data/MLOps)**

- DVC-tracked proto\_pack (10–20 CTs); dataset card.
- FastAPI + queue resilient to kill/restart; metrics scripts produce CSVs.
- README with one-command run; environment pins.

**Rahul (Frontend)**

- Ingest → Analyze → Results works offline; overlays toggle smoothly.
  - Bilingual toggle; all downloads work.
  - Error states friendly; show queued FHIR JSONs.
- 

**Quick tips / pitfalls per role**

- **Dinesh:** If UNet is slow, reduce cube size to 48 or stride to 2; keep batch small (fits 8 GB). Accept coarse masks for proto; quality later.
  - **Kowshik:** Use a Devanagari font (e.g., Noto Sans Devanagari) so Hindi renders in PDFs. Keep Q&A strictly **schema-grounded** (no free-text LLM yet).
  - **Teja:** Don't load entire volume into RAM for all steps; batch slices; log per-stage time. Persist job state to DB so resume works.
  - **Rahul:** Precache overlays (PNG) so toggles are instant. Window/level on the client with numpy → PIL conversions.
- 

If you want, I can also drop **starter files** (skeleton code for each of the modules above), plus a **sample findings.json**, **FHIR JSON**, and **report templates (EN/HI)** so you can kick off implementation today.